Five Minute Video Demo Script for Social Pulse

Show a working Social Media Analytics Dashboard built for Week 1, Path 3, using Streamlit and vibe coding with Codex. The timed walkthrough runs from 0:00 to 5:00 and covers CSV upload, visual analysis, shared filters, AI summaries, natural-language questions, and the development process. Read the quoted narration aloud; the screen actions and recording notes are presenter cues.

Prepare before recording

1. Open the running dashboard at **http://127.0.0.1:8501/** in a wide browser window. Begin on **Overview**, with **Sample dataset** selected, all categories included, search blank, and **Date range → All data**. Keep the sidebar visible.
2. Have social_media_outdoor_sample.csv ready in the file picker. This is the second scenario: 1,080 fictional Cedar & Summit posts from January through June 2026. Both CSVs have the same required 18 columns. The original has 720 fictional posts from July through September 2026.
3. Configure your API key privately in **AI assistant → AI connection settings**. Keep that panel collapsed and caption sharing off during recording. Rehearse a summary and the question below using the outdoor dataset; verify the returned calculations. Then clear AI history and return to the opening state. The build's automated AI tests use mocked responses, so this rehearsal is the live integration check.
4. Open the comprehensive project report beside the browser, ready at its actual prompts and development iterations. Alternatively, use the actual development conversation. If using **PROJECT_NOTES.md**, show its **Actual starting request** and **AI interactions added in a follow-up iteration** sections; the six reusable stage prompts are examples, not six messages you submitted.
5. Copy this question into a separate note so you can paste it quickly at 3:20:

> Using this synthetic dataset, which paid campaign has the highest ROAS? Show spend, revenue, ROAS, and the number of paid posts.

6. Rehearse at a natural pace of roughly 125-135 words per minute. The main narration leaves room for clicks and AI response time. Use the fixed timestamps as transition cues. If an AI call fails or runs too long, use the replacement below and finish by 5:00.

CSV upload and visual exploration

1 Introduce the project

Time: 0:00-0:25 · 25 seconds

On screen: Show the title, Overview KPI cards, and sidebar. Point to the sample source and 720-post count.

NARRATION

> “This is Social Pulse, my Social Media Analytics Dashboard for Week One, Path Three. I built it with Streamlit and Codex to turn a CSV export into charts, filters, and useful analysis. The starting sample contains seven hundred and twenty fictional posts. All data in this demo is synthetic.”

2 Upload a different CSV

Time: 0:25-1:05 · 40 seconds

On screen: Select **Dataset → Upload CSV**, click **Choose a CSV**, and select **social_media_outdoor_sample.csv**. After loading, set **Date range → All data**. Point to the filename, January-June dates, and **1,080 of 1,080 posts**. Briefly show the new platform choices, including Pinterest and Reddit.

NARRATION

> “I can switch from the built-in sample to a compatible CSV. Here I’m uploading the second dataset, created for a fictional outdoor retailer called Cedar and Summit. It contains one thousand and eighty posts across six months, with different campaigns, captions, and platform coverage. I select All data to include the whole file. The uploader checks required columns, dates, and numeric values before the dashboard displays the results.”

3 Explore charts and the post table

Time: 1:05-1:45 · 40 seconds

On screen: Point to the Overview KPIs and trend chart. Change **Group dates** from **Weekly** to **Monthly**. Open **Content performance**, hover over a format bar and a scatterplot point, then briefly scroll to **Post explorer**. Change **Rank posts by → Clicks** and **Posts to show → 10**.

NARRATION

> “The overview shows impressions, engagement actions, clicks, and attributed new followers. Engagement rate uses total actions divided by summed post reach, so it is weighted by reach. I can group the trend by month, then compare content formats. Each scatterplot dot represents one post, with bubble size showing clicks. Below, I can rank individual posts by clicks and inspect the top ten in the table.”

Filters and analytical workflows

4 Apply shared filters

Time: 1:45-2:15 · 30 seconds

On screen: Keep **Content performance** visible. Set **Date range → Last 30 days**. Remove the other platform chips so **Platforms** contains only **Instagram**; leave the other category filters selected. Point to the filtered post count and changed charts. Click **Reset filters**, then set **Date range → All data** again.

NARRATION

> “The sidebar filters every view. I’ll choose the dataset’s latest thirty days and keep only Instagram. The charts and table update together. Formats, campaigns, traffic type, and keywords are also available. I’ll reset the filters and select All data again before asking AI about the complete file.”

5 Show campaign analysis and a timing experiment

Time: 2:15-2:40 · 25 seconds

On screen: Open **Campaigns** and point to **Paid ROAS** and the paid spend versus attributed revenue chart. Open **Analysis lab** and hover over a populated weekday/hour heatmap cell. Keep **Minimum posts per time slot** at 3.

NARRATION

> “Campaigns compares paid spend and attributed revenue. ROAS uses paid posts and measures return on ad spend, not profit. The Analysis lab’s timing heatmap hides slots with too few posts, helping me choose a publishing-time hypothesis to test with similar content.”

6 Generate an AI summary

Time: 2:40-3:20 · 40 seconds

On screen: Open **AI assistant**. Confirm the scope shows **1,080 posts** and keep connection settings collapsed. Click **Generate AI summary** immediately, then narrate while it runs. When the response appears, point to one numerical observation and expand **View supporting calculations**. Do not read the entire response aloud.

NARRATION

> “The AI assistant works with the current dashboard filters. I’ll generate a summary of this synthetic dataset. The model can request validated calculations that the application performs locally, then explain the results. When the summary appears, I can inspect the supporting figures. That lets me check a numerical observation against the data before treating a suggested experiment as useful.”

Natural language questions and vibe coding

7 Ask a question in natural language

Time: 3:20-4:00 · 40 seconds

On screen: Paste the prepared question into **Your question about the CSV** and click **Ask AI** immediately. Narrate while it runs. When the answer appears, expand its **View supporting calculations** and point to the campaign, paid post count, spend, revenue, and ROAS. Use the actual returned values; no particular answer needs to be pre-scripted.

NARRATION

> “Now I’ll ask: ‘Using this synthetic dataset, which paid campaign has the highest ROAS? Show spend, revenue, ROAS, and the number of paid posts.’ The answer includes supporting calculations, so I can compare the explanation with the campaign totals. Questions can narrow the current filters. To include excluded dates or platforms, I need to change the sidebar first.”

If comfortably ahead of time: Show that a follow-up can ask **“Show its top five paid posts by revenue.”** Submit it only if the rehearsal shows both calls fit within this section; otherwise demonstrate the single question and its evidence.

8 Explain the vibe coding workflow

Time: 4:00-4:50 · 50 seconds

On screen: Switch to the project report's actual prompts and iterations, or the real development conversation. Point to the initial build request, the AI follow-up, and the request for different sample data. Briefly show the refinement and validation notes.

NARRATION

> “I used vibe coding by describing the outcome to Codex, then reviewing and refining the implementation. My first prompt specified Streamlit, CSV uploads, charts, filters, and UX improvements. Next, I asked for AI summaries and natural-language questions. A third prompt requested completely different sample data, producing this outdoor retailer scenario. Refinements addressed currency formatting, consistent filters, and chat updates without an extra rerun. The build passed forty-three automated checks, with AI responses mocked. My main learning was that clear prompts speed up development, but reliable analysis still depends on explicit metric definitions, validation, and checking the evidence.”

9 Close the demo

Time: 4:50-5:00 · 10 seconds

On screen: Return to Overview. Point to **Export filtered posts** and finish on the dashboard title. Stop recording at 5:00.

NARRATION

> “Social Pulse combines CSV uploads, charts, filters, and AI questions to help me plan the next content experiment.”

Replacement if live AI is unavailable

Use these two blocks in place of sections 6 and 7, keeping the same 2:40-4:00 time window. Show the actual setup or error state. This demonstrates the AI controls and local alternative; it does not demonstrate a successful live AI response. If a successful AI response is required for submission, resolve access and record those sections again.

At 2:40-3:20: Show **AI assistant**, **Generate AI summary**, and **Your question about the CSV**. If a request fails, point to the displayed message without repeated retries. Keep credentials hidden.

> “This tab provides AI summaries and natural-language questions for the currently filtered CSV. A live response requires a configured API key and an available service. That connection is unavailable in this recording, so I’m showing the controls and their scope. Once connected, the model requests local calculations and returns an explanation with evidence. The rest of the dashboard continues to work.”

At 3:20-4:00: Open **Overview → What the data shows**, then **Analysis lab → Choose an analysis question → Which paid campaign has the highest ROAS?** Point to the computed result and evidence. Leave **Minimum posts per time slot** at 3; the preset also requires at least three paid posts per campaign.

> “These summaries and preset questions are calculated locally, rather than generated by AI. Here I can still ask which paid campaign has the highest return on ad spend and review the underlying numbers. This preset requires at least three paid posts per campaign. It keeps the workflow useful while the live AI connection is unavailable.”

Recording notes

- Keep visible filter scope and post counts in frame so viewers know which data each result describes. **Last 30 days** is relative to the file's latest date, not today's date. **Reset filters** returns to that default; it does not select All data.
- Keep at least one value in every category filter. Clearing all chips produces no matching posts. Keep both Paid and Organic selected for the AI demonstration so the model can narrow to paid campaigns itself.
- Make the summary request and one question the priority. Skip the optional follow-up if it threatens the five-minute limit. If a response has not appeared by the section boundary, state that it is still pending and move on; do not narrate an answer that is absent.
- Use the actual prompts as development evidence. The original starting request described the six-step workflow in one message; avoid presenting the reusable stage prompts as messages you sent.
- Refer to the uploaded outdoor CSV as synthetic in narration and in the AI question. Its generated patterns are practice material, not findings about real social platforms.
- Rehearse the upload picker, scrolling, and switch to the prompt evidence before recording. If trimming pauses afterward, label the recording as edited. Aim to finish slightly before 5:00 if your submission platform enforces a strict duration limit.